



DNN Accelerator Design Optimization and NoC-based DNN

Kun-Chih (Jimmy) Chen 陳坤志

kcchen@nycu.edu.tw

Institute of Electronics,


National Yang Ming Chiao Tung University



Complexity in ANN (1/2)

- ❖ One layer of fully-connected (FC) ANN layer can be represented by an $n \times m$ matrix containing weights in links from m input neurons to n output neurons
 - ❖ e.g., 3 input neurons, 4 output neurons \Rightarrow 4×3 weight matrix plus 4×1 vector of biases
- ❖ In general, there are hundreds (thousands) input/output neurons, and several FC ANN layers (for classification)
 - ❖ FC-1 in VGG-16: input neurons: $7 \times 7 \times 512 = 25088$; output neurons: 4090 \Rightarrow 102,760,000 parameters
 - ❖ FC-2 in VGG-16: input neurons: 4096; output neurons: 4090 \Rightarrow 16,777,216 parameters
 - ❖ FC3 in VGG-16: input neurons: 4090; output neurons: 1000 \Rightarrow 4,096,000 parameters
 - ❖ Total number of parameters in VGG-16: 138M \Rightarrow about 90% of parameters in FC layers


Complexity in ANN (2/2)



total # of weights (links): $3 \times 4 + 4 \times 4 + 4 \times 1 = 32$

total # of biases: $4 + 4 + 1 = 9$

total # of operations (MAC): $3 \times 4 + 4 \times 4 + 4 \times 1 = 32$



$$\begin{bmatrix} y_1 \\ y_2 \\ y_3 \\ y_4 \end{bmatrix} = \underbrace{\begin{bmatrix} w_{11} & w_{12} & w_{13} \\ w_{21} & w_{22} & w_{23} \\ w_{31} & w_{32} & w_{33} \\ w_{41} & w_{42} & w_{43} \end{bmatrix}}_M \begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix} + \begin{bmatrix} b_1 \\ b_2 \\ b_3 \\ b_4 \end{bmatrix}$$


total # of weights (links): $3 \times 4 = 12$

total # of biases: 4

total # of operations (MAC): $3 \times 4 = 12$

Reduced Computations

- ❖ $MM \times M \rightarrow MN \times NN \times NM \times M$
- ❖ # of multiplications: $M^2 \rightarrow 2MN + N^2$
 - ❖ i.e., $M=100, N=20$
 - ❖ $10000 \rightarrow 4400$ (56% reduction)



Low-rank approximation: truncated SVD

- ❖ Fully connected (FC) layers can be represented by matrix-vector product
 - ❖ Matrix containing weights can be decomposed using SVD
 - ❖ Reduce # of weights by selecting only the largest singular values

$$\begin{array}{ccccc}
 \begin{array}{c} m \\ \boxed{} \\ n \end{array} & = & \begin{array}{c} r \\ \boxed{\begin{array}{c} | | | | | \\ U_1 U_2 U_3 \dots \end{array}} \\ n \end{array} & \begin{array}{c} r \\ \boxed{\begin{array}{ccc} S_1 & & 0 \\ & S_2 & \\ 0 & & \ddots \\ & & & S_r \end{array}} \\ r \end{array} & \begin{array}{c} m \\ \boxed{\begin{array}{c} \text{---} V_1 \text{---} \\ \text{---} V_2 \text{---} \\ \text{---} V_3 \text{---} \\ \vdots \end{array}} \\ r \end{array} \\
 X & & U & S & V^T \\
 \\
 \begin{array}{c} m \\ \boxed{\phantom{\hat{X}}} \\ n \end{array} & = & \begin{array}{c} k \\ \boxed{\begin{array}{c} | \phantom{\hat{U}_1} | \phantom{\hat{U}_2} | \phantom{\hat{U}_3} | | \\ \hat{U}_1 \hat{U}_2 \hat{U}_3 \dots \end{array}} \\ n \end{array} & \begin{array}{c} k \\ \boxed{\begin{array}{ccc} \hat{S}_1 & & 0 \\ & \hat{S}_2 & \\ 0 & & \ddots \\ & & & \hat{S}_k \end{array}} \\ k \end{array} & \begin{array}{c} m \\ \boxed{\begin{array}{c} \text{---} \hat{V}_1 \text{---} \\ \text{---} \hat{V}_2 \text{---} \\ \text{---} \hat{V}_3 \text{---} \\ \vdots \end{array}} \\ k \end{array} \\
 \hat{X} & & \hat{U} & \hat{S} & \hat{V}^T
 \end{array}$$


Neural Network Model Compression

- ❖ Selecting first 500 (out of 4096) largest singular values for Fully-Connected layer FC6 with 14x14x128 input neurons and 4096 output neurons
- ❖ 2/3 of total weights are cut off after SVD decomposition

$$f^{out} = W f^{in} + b \Rightarrow W \approx U_d S_d V_d \quad \begin{matrix} U_d \in \mathbb{R}^{n_{out} \times d}, V_d \in \mathbb{R}^{d \times n_{in}} \\ S_d \in \mathbb{R}^{d \times d} \end{matrix}$$

$$O(n_{in}n_{out}) \Rightarrow O(dn_{in} + dn_{out}) \quad d \ll n_{in}, n_{out}$$

Network	FC6	# of total weights	# of operations	Top-5 accuracy
VGG-16	25088×4096	138.36M	30.94G	88.00%
VGG-16-SVD	$25088 \times 500 + 500 \times 4096$	50.18M	30.76G	87.96%



VGG16

Why do we need DNN model compression?

❖ Efficient deployment

- ❖ Efficient deployment of neural networks on devices with limited resources
- ❖ Faster inference and reduced energy consumption

❖ Reduce cost

- ❖ Smaller models can reduce the cost of hardware and infrastructure needed to train and deploy neural networks

❖ Energy efficiency

- ❖ Model compression can reduce the power consumption of DNN models, which is essential for power resource-limited edge devices.

Techniques of DNN Model Compression

❖ Pruning

- ❖ Removing the least important connections, neurons, or filters from the network
→ Making it smaller and faster.

❖ Quantization

- ❖ Reducing the precision of the weights and activations in the network. → Reducing the memory requirements and computational complexity of the model.

❖ Weight Sharing

- ❖ Sharing weights between multiple layers or filters in the network. → Reducing the number of unique weights in the model, making it smaller and faster.

❖ Depthwise Separable Convolution

- ❖ Each channel is convolved with its own set of filters, rather than all channels being convolved with the same filter → Reducing the number of parameters needed in convolution layers.



Pruning



Pruning Neural Network

- ❖ Removing weights of small magnitudes recursively
 - ❖ Remove connections with weight below a per-layer adjustable threshold
 - threshold chosen as a quantity parameter multiplied by standard deviation of a layer's weights
 - ❖ Retrain dropout ratio should be smaller




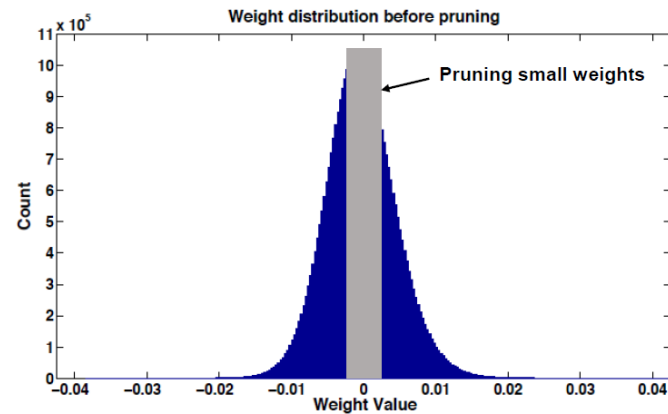


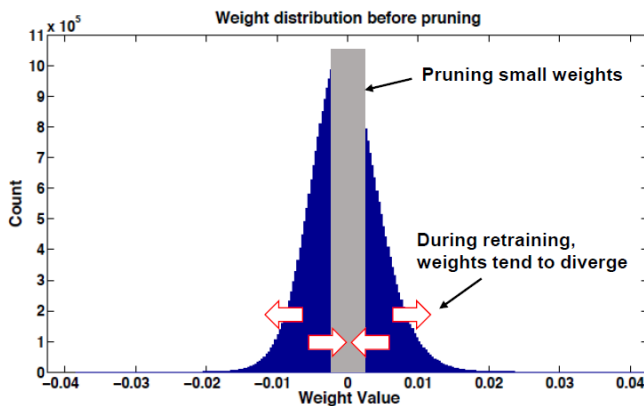
Illustration of Iterative Pruning-Retraining




weight distribution
before pruning



small weights are set to zero
(pruned)



during training, weights
are re-distributed



10X smaller scale

after a step of training, new
small weights are obtained

Accuracy, Compression Ratio and # Bits (Quantization Levels)

Network	Top-1 Error	Top-5 Error	Parameters	Compress Rate
LeNet-300-100 Ref	1.64%	-	1070 KB	
LeNet-300-100 Compressed	1.58%	-	27 KB	40×
LeNet-5 Ref	0.80%	-	1720 KB	
LeNet-5 Compressed	0.74%	-	44 KB	39×
AlexNet Ref	42.78%	19.73%	240 MB	
AlexNet Compressed	42.78%	19.70%	6.9 MB	35×
VGG-16 Ref	31.50%	11.32%	552 MB	
VGG-16 Compressed	31.17%	10.91%	11.3 MB	49×


AlexNet	#CONV bits / #FC bits	Top-1 Error	Top-5 Error	Top-1 Error Increase	Top-5 Error Increase
	32bits / 32bits	42.78%	19.73%	-	-
	8 bits / 5 bits	42.78%	19.70%	0.00%	-0.03%
	8 bits / 4 bits	42.79%	19.73%	0.01%	0.00%
	4 bits / 2 bits	44.77%	22.33%	1.99%	2.60%

AlexNet Pruning

- ❖ # of parameters from 61M to 6.7M (11%)
- ❖ # of operations from 1.5Gop to 0.45Gop (30%)
- ❖ More FC-layer weights removed cp. Conv layers
- ❖ Retrain with 1/100 of the original network's original learning rate
 - ❖ cp. with 1/10 learning rate in LeNet

Layer	Weights	FLOP	Act%	Weights%	FLOP%
conv1	35K	211M	88%	84%	84%
conv2	307K	448M	52%	38%	33%
conv3	885K	299M	37%	35%	18%
conv4	663K	224M	40%	37%	14%
conv5	442K	150M	34%	37%	14%
fc1	38M	75M	36%	9%	3%
fc2	17M	34M	40%	9%	3%
fc3	4M	8M	100%	25%	10%
Total	61M	1.5B	54%	11%	30%

some colors missing below



VGG-16 Pruning

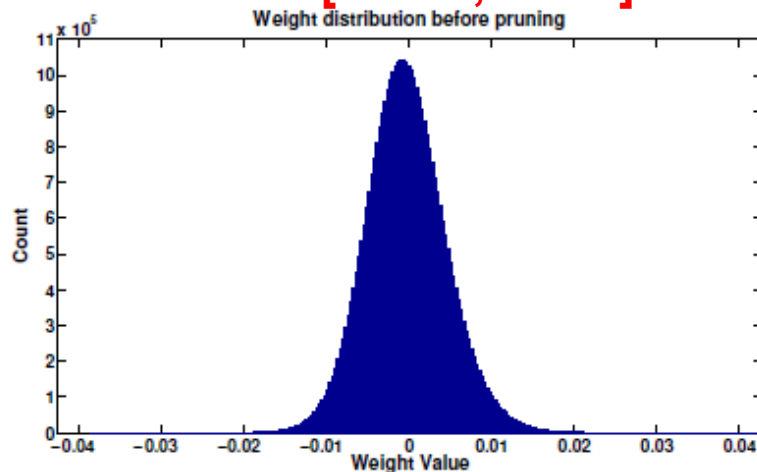
- ❖ # of parameters from 138M to 10.35M (7.5%)
- ❖ Operation count from 30.9 GOP to 6.5 GOP (21%)
- ❖ After pruning of AlexNet and VGG, weights can be stored on chip

Layer	Weights	FLOP	Act%	Weights%	FLOP%
conv1_1	2K	0.2B	53%	58%	58%
conv1_2	37K	3.7B	89%	22%	12%
conv2_1	74K	1.8B	80%	34%	30%
conv2_2	148K	3.7B	81%	36%	29%
conv3_1	295K	1.8B	68%	53%	43%
conv3_2	590K	3.7B	70%	24%	16%
conv3_3	590K	3.7B	64%	42%	29%
conv4_1	1M	1.8B	51%	32%	21%
conv4_2	2M	3.7B	45%	27%	14%
conv4_3	2M	3.7B	34%	34%	15%
conv5_1	2M	925M	32%	35%	12%
conv5_2	2M	925M	29%	29%	9%
conv5_3	2M	925M	19%	36%	11%
fc6	103M	206M	38%	4%	1%
fc7	17M	34M	42%	4%	2%
fc8	4M	8M	100%	23%	9%
total	138M	30.9B	64%	7.5%	21%

Comparison (AlexNet)


Network	Top-1 Error	Top-5 Error	Parameters	Compression Rate
Baseline Caffemodel [26]	42.78%	19.73%	61.0M	1×
Data-free pruning [28]	44.40%	-	39.6M	1.5×
Fastfood-32-AD [29]	41.93%	-	32.8M	2×
Fastfood-16-AD [29]	42.90%	-	16.4M	3.7×
Collins & Kohli [30]	44.40%	-	15.2M	4×
Naive Cut	47.18%	23.23%	13.8M	4.4×
SVD [12]	44.02%	20.56%	11.9M	5×
Network Pruning	42.77%	19.67%	6.7M	9×

most in $[-0.015, 0.015]$



Weight distribution before parameter pruning

most in $[-0.025, 0.025]$




Weight distribution after parameter pruning

Why Increase Sparsity?


- ❖ Reduce number of MACs
 - ❖ Anything multiplied by zero is zero → avoid performing unnecessary MACs
 - ❖ Reduce energy consumption and latency
- ❖ Reduce data movement
 - ❖ If one of the inputs to MAC is zero, can avoid reading the other input
 - ❖ Compress data by only sending non-zero values
- ❖ CPU/GPU libraries typically only support really high sparsity (> 99%) due to the overhead
 - ❖ Sparsity for DNNs typically much lower → need specialized hardware

Sparsity in Activation Data

Many **zeros** in output fmaps after ReLU



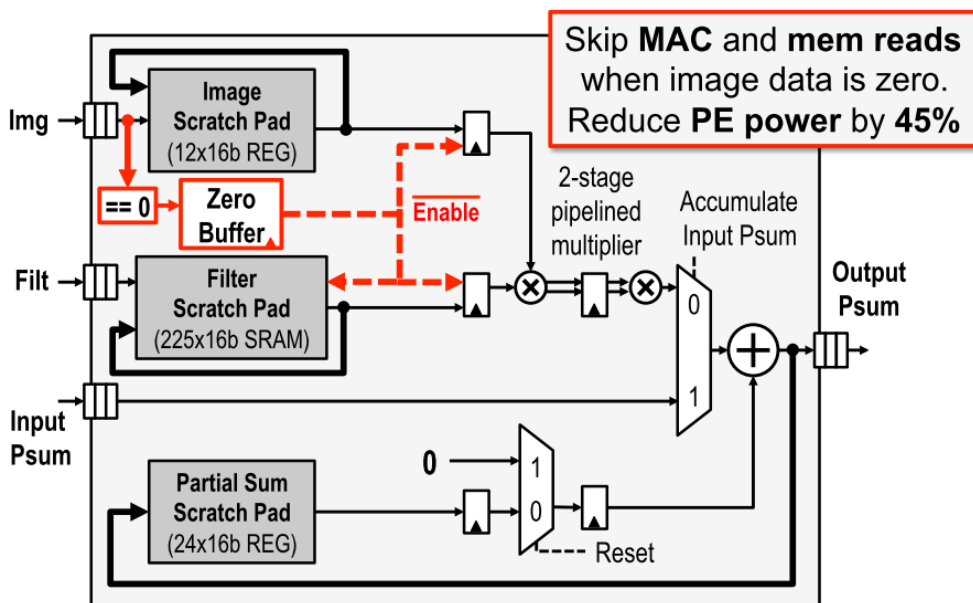
■ # of activations ■ # of non-zero activations



[Chen, ISSCC 2016]


Data Gating / Zero Skipping

Gate operations (reduce power consumption)



Eyeriss [Chen, ISSCC 2016]

Skip operations (increase throughput)



Cnvlutin [Albericio, ISCA 2016]

Unstructured or Structured Sparsity




Image Source: [**Mao**, CVPR Workshop 2017]

❖ Benefits

- ❖ Increase coarseness \rightarrow more structure in sparsity (easier for hardware)
- ❖ Less signaling overhead for location of zeros \rightarrow better compression

Design Considerations for Sparsity

❖ Impact on accuracy

- ❖ Must consider difficulty of dataset, task, and DNN model
 - e.g., AlexNet and VGG known to be over parameterized and thus easy to prune weights; does method work on efficient DNN models?

❖ Does hardware cost exceed benefits?

- ❖ Need extra hardware to identify sparsity
 - e.g., Additional logic to identify non-zeros and store non-zero locations
- ❖ Accounting for sparsity in both weights and activations is challenging
 - Need to compute intersection of two data streams rather than find next non-zero in one
- ❖ Granularity impacts hardware overhead as well as accuracy
 - e.g., Fine-grained or coarse-grained (structured) sparsity
- ❖ Compressed data will be variable length
 - Reduced flexibility in access order → random access will have significant overhead




Quantization and Weight Sharing



Quantization and Weight Sharing

- Quantize to fixed number of distinct values at no accuracy loss
- AlexNet conv layers quantized using 8 bits (256 16-bit weights) results in zero accuracy loss



Quantization and Weight Sharing

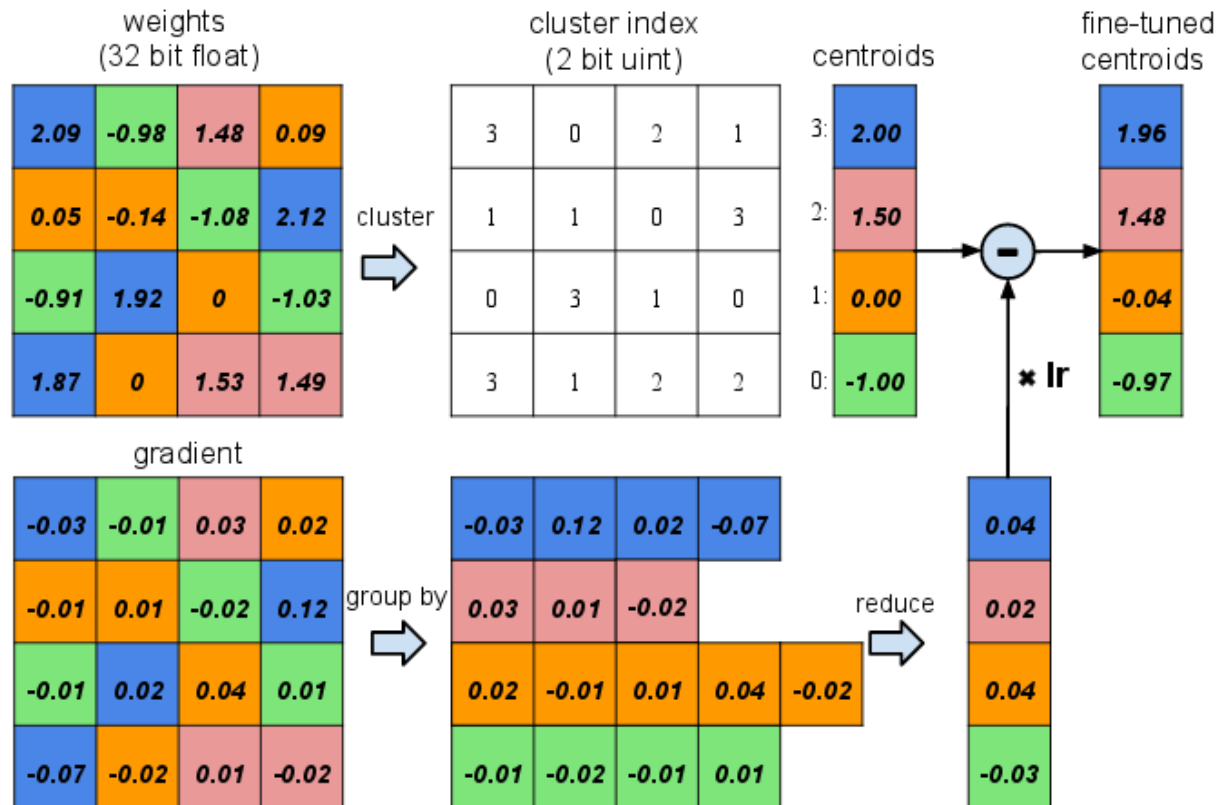
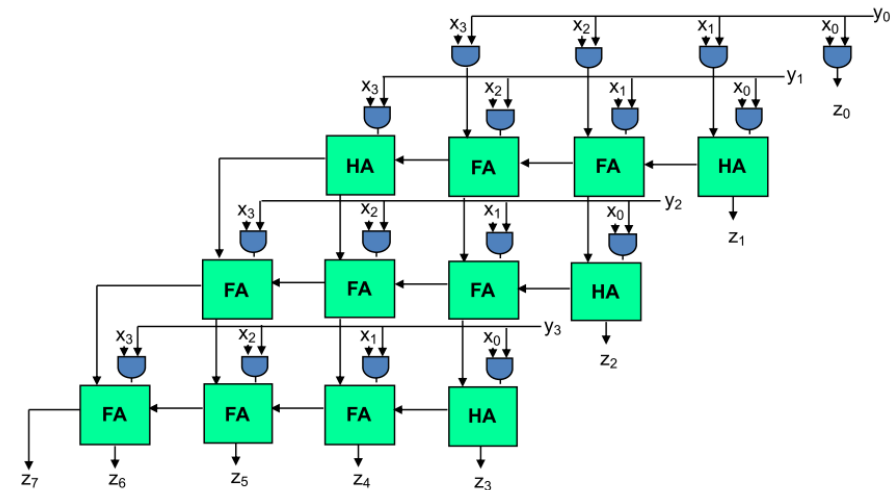


Figure 3: Weight sharing by scalar quantization (top) and centroids fine-tuning (bottom).

Why Reduce Precision (i.e., Reduce Bit Width)?


- ❖ Reduce data movement and storage cost for inputs and outputs of MAC
 - ❖ Smaller memory \rightarrow lower energy
- ❖ Reduce cost of MAC
 - ❖ Cost of multiply increases with bit width (n) \rightarrow energy and area by $O(n^2)$; delay by $O(n)$

					x_3	x_2	x_1	x_0	Multiplicand
					y_3	y_2	y_1	y_0	Multiplier
					x_3y_0	x_2y_0	x_1y_0	x_0y_0	
					x_3y_1	x_2y_1	x_1y_1	x_0y_1	Partial Product
					x_3y_2	x_2y_2	x_1y_2	x_0y_2	
					x_3y_3	x_2y_3	x_1y_3	x_0y_3	
+	z_7	z_6	z_5	z_4	z_3	z_2	z_1	z_0	Result



Note: Bit width for multiplication and accumulation in a MAC are different


Impact of Reduced Precision on Energy & Area



What Determines Bit Width? (1/4)

- ❖ Number of unique values
 - ❖ e.g., **M-bits** to represent 2^M values
- ❖ Dynamic range of values
 - ❖ e.g., **E-bits** to scale values by $2^{(E-127)}$
- ❖ Signed or unsigned values
 - ❖ e.g., signed requires one extra bit (S)
- ❖ **Total bits = S+E+M**
- ❖ **Floating point (FP)** allows range to change for each value (E-bits)
- ❖ **Fixed point (Int)** has fixed range
- ❖ Default CPU/GPU is 32-bit float (**FP32**)

Common Numerical Representations



What Determines Bit Width? (2/4)

- ❖ For accuracy, require sufficient precision to represent different data types
 - ❖ For inference: weights, activations, and partial sums
 - ❖ For training: weights, activations, partial sums, gradients, and weight update
 - ❖ Required precision can vary across data types
 - ❖ Referred to as mixed precision

What Determines Bit Width? (3/4)

- ❖ **Reduce number of unique values (M-bits, a.k.a. mantissa)**
 - ❖ **Default:** Uniform quantization (values are equally spaced out)
 - ❖ Non-uniform quantization (spacing can be computed, e.g., logarithmic, or with look-up-table)
 - ❖ Fewer unique values can make transforms and compression more effective

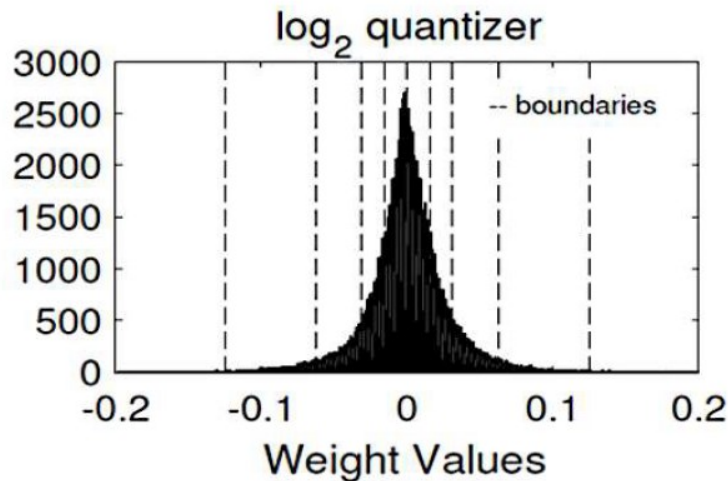



Image Source: [Lee, ICASSP 2017]

What Determines Bit Width? (4/4)

- ❖ **Reduce number of unique values (M-bits, a.k.a. mantissa)**
 - ❖ **Default:** Uniform quantization (values are equally spaced out)
 - ❖ Non-uniform quantization (spacing can be computed, e.g., logarithmic, or with look-up-table)
 - ❖ Fewer unique values can make transforms and compression more effective
- ❖ **Reduce dynamic range (E-bits, a.k.a., exponent)**
 - ❖ If possible, fix range (i.e., used fixed point, E=0)
 - ❖ Share range across group of values (e.g., weights for a layer or channel)
- ❖ Tradeoff between number of bits allocated to **M-bits** and **E-bits**

fp16 (S=1, E=5, M=10)  range: $\sim 5.9e^{-8}$ to $\sim 6.5e^4$

bfloat16 (S=1, E=8, M=7)  range: $\sim 1e^{-38}$ to $\sim 3e^{38}$

Design Considerations for Reduced Precision

❖ Impact on accuracy

- ❖ Must consider difficulty of dataset, task, and DNN model
 - e.g., Easy to reduce precision for an easy task (e.g., digit classification); does method work for a more difficult task?

❖ Does hardware cost exceed benefits?

- ❖ Need extra hardware to support variable precision
 - e.g., Additional shift-and-add logic and registers for variable precision
- ❖ Granularity impacts hardware overhead as well as accuracy
 - e.g., More overhead to support (1b, 2b, 3b ... 16b) than (2b, 4b, 8b, 16b)

❖ Evaluation

- ❖ Use 8-bit for inference and 16-bit float for training for baseline
- ❖ 32-bit float is a weak baseline




Depthwise Separable Convolution (DSC)




Depthwise Separable Convolution (1/7)

❖ Original Convolution



Depthwise Separable Convolution (2/7)


- ❖ Depthwise separable convolution is proposed to reduce the amount of computation without affecting the output structure.
- ❖ It can be split into two parts:
 - ❖ Depthwise Convolution
 - ❖ Pointwise Convolution



Depthwise Separable Convolution (3/7)


- ❖ Standard convolution

- ❖ M input channels, N output channels
- ❖ $D_K \times D_K$ filters




- ❖ Depthwise convolution (DWC)

- ❖ Convolution of one input channel, generating one output channel
- ❖ # of output channels = # of input channels



- ❖ Pointwise convolution (PWC)

- ❖ 1×1 filters combining M input channels, generating N output channels




- ❖ DSC = DWC + PWC

Depthwise Separable Convolution (4/7)


❖ Depthwise Convolution

Input data



Depthwise Separable Convolution (5/7)

❖ Pointwise Convolution



Depthwise Separable Convolution (6/7)

- ❖ # Original Convolution operations can be approximated to:

$$H_{out} \times W_{out} \times C \times K \times K \times N$$

- ❖ # Depthwise Separable Convolution operations can be approximated to:

$$\boxed{H_{out} \times W_{out} \times C \times K \times K} + \boxed{H_{out} \times W_{out} \times C \times N}$$

Depthwise Pointwise

Op. reduction =

$$\frac{H_{out} \times W_{out} \times C \times K \times K + H_{out} \times W_{out} \times C \times N}{H_{out} \times W_{out} \times C \times K \times K \times N}$$
$$= \frac{1}{K \times K} + \frac{1}{N}$$

Depthwise Separable Convolution (7/7)

$$\begin{aligned}
 & \frac{H_{out} \times W_{out} \times C \times K \times K + H_{out} \times W_{out} \times C \times N}{H_{out} \times W_{out} \times C \times K \times K \times N} \\
 = & \frac{1}{K \times K} + \frac{1}{N}
 \end{aligned}$$

```

=====
Total params: 138,357,544
Trainable params: 138,357,544
Non-trainable params: 0
    
```


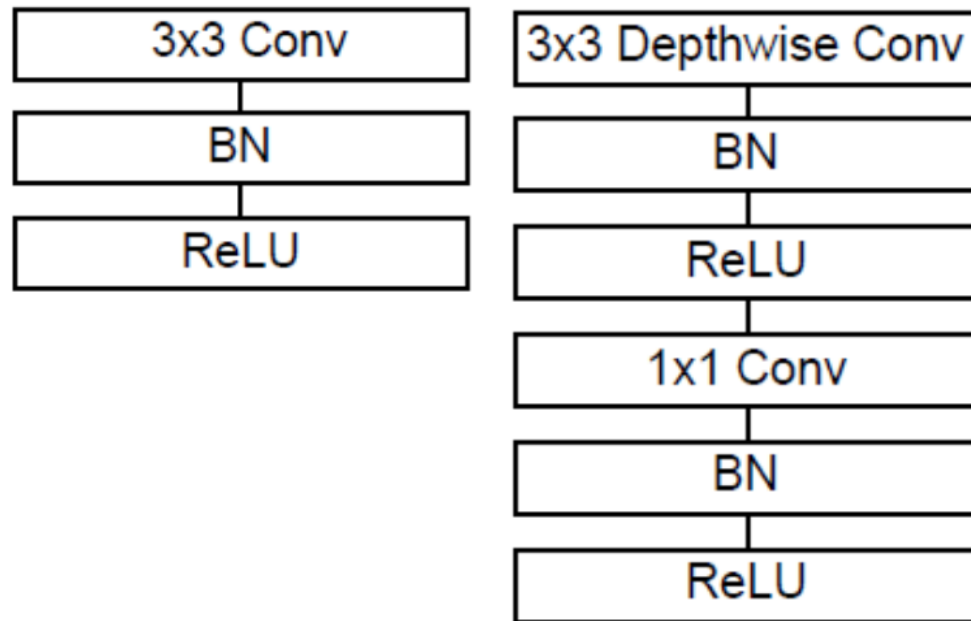
```

=====
Total params: 4,253,864
Trainable params: 4,231,976
Non-trainable params: 21,888
    
```

Model	ImageNet Accuracy	Million Mult-Adds	Million Parameters
1.0 MobileNet-224	70.6%	569	4.2
GoogleNet	69.8%	1550	6.8
VGG 16	71.5%	15300	138

Light CNN Models: MobileNet

❖ Standard 3D Convolution = Depthwise Convolution + 1x1 Convolution



MobileNet

❖ Width multiplier

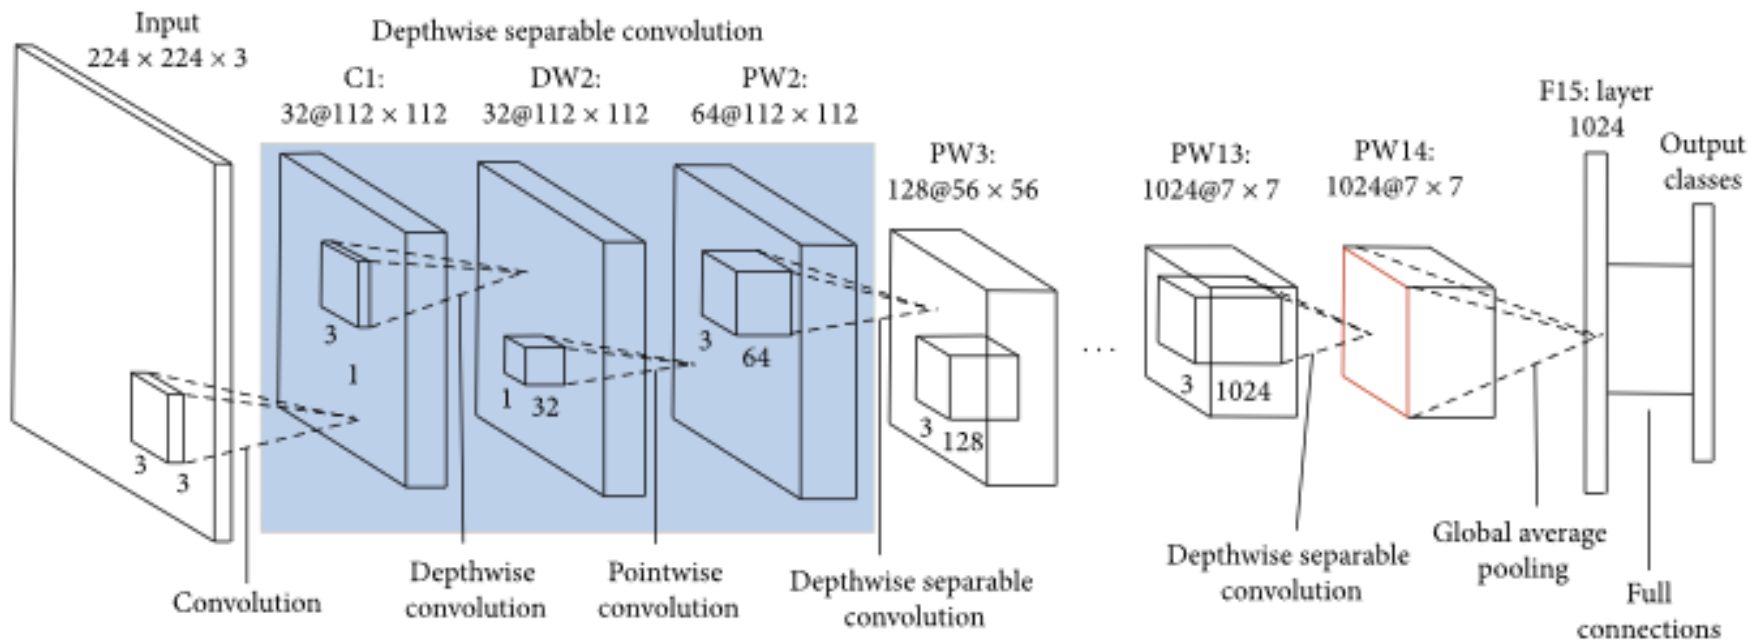
- ❖ Reduced channel number

❖ Resolution multiplier

- ❖ Reduced channel size

Table 8. MobileNet Comparison to Popular Models

Model	ImageNet Accuracy	Million Mult-Adds	Million Parameters
1.0MobileNet-224	70.6%	569	4.2
GoogleNet	69.8%	1550	6.8
VGG 16	71.5%	15300	138







Flexible and Scalability



Many Efficient DNN Design Approaches



Reduce Precision

32-bit float 1 0 1 0 0 1 0 1 0 0 0 0 0 0 0 0 0 1 0 1 0 0 0 0 0 0 0 0 0 0 0 1 0 0

8-bit fixed 0 1 1 0 0 1 1 0


Binary 0

No guarantee that DNN algorithm designer will use a given approach.
Need flexible DNN processor!

[Chen, SysML 2018]

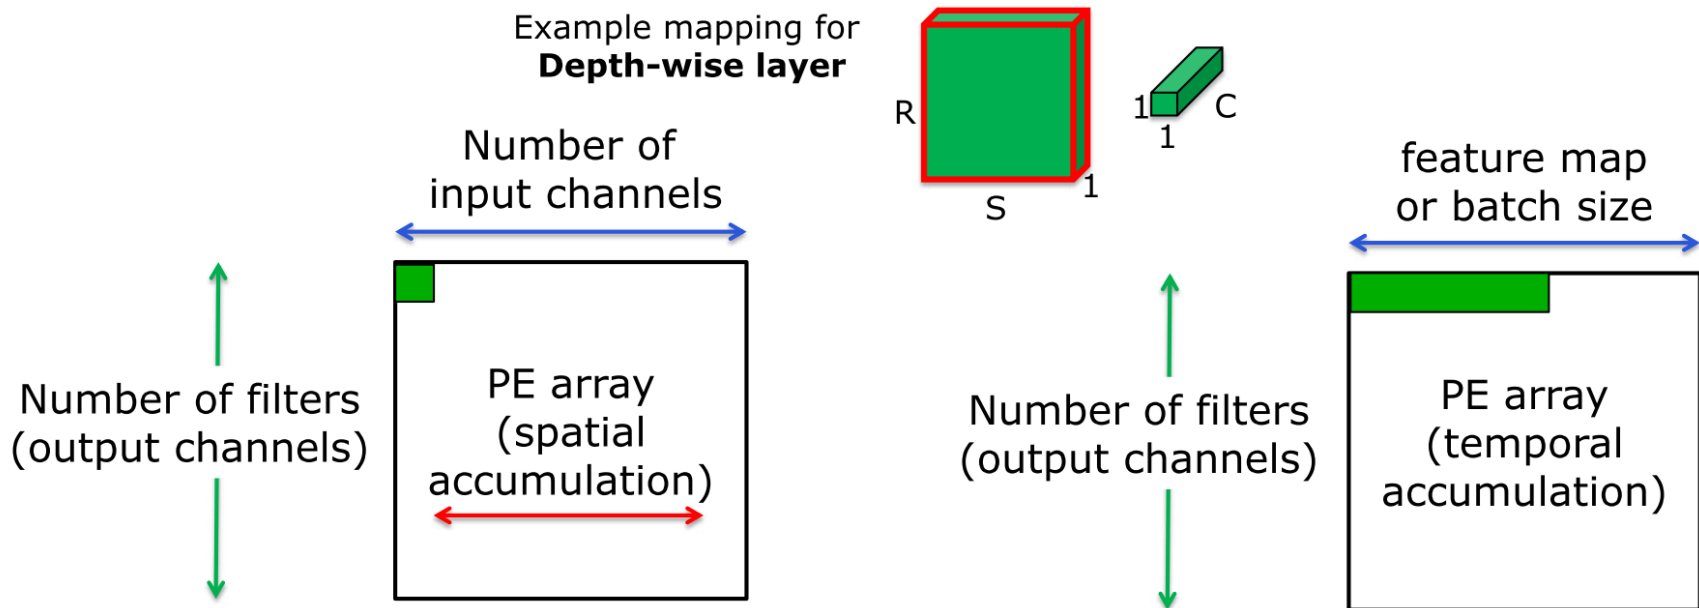
Limitations of Existing DNN Processors (1/2)

- ❖ Specialized DNN processors often rely on certain properties of the DNN model in order to achieve high energy-efficiency
- ❖ Example: Reduce memory access by amortizing across PE array




Limitations of Existing DNN Processors (2/2)

- ❖ Reuse depends on # of channels, feature map/batch size
 - ❖ Not efficient across all DNN models (e.g., efficient network architectures)



Need Flexible Dataflow

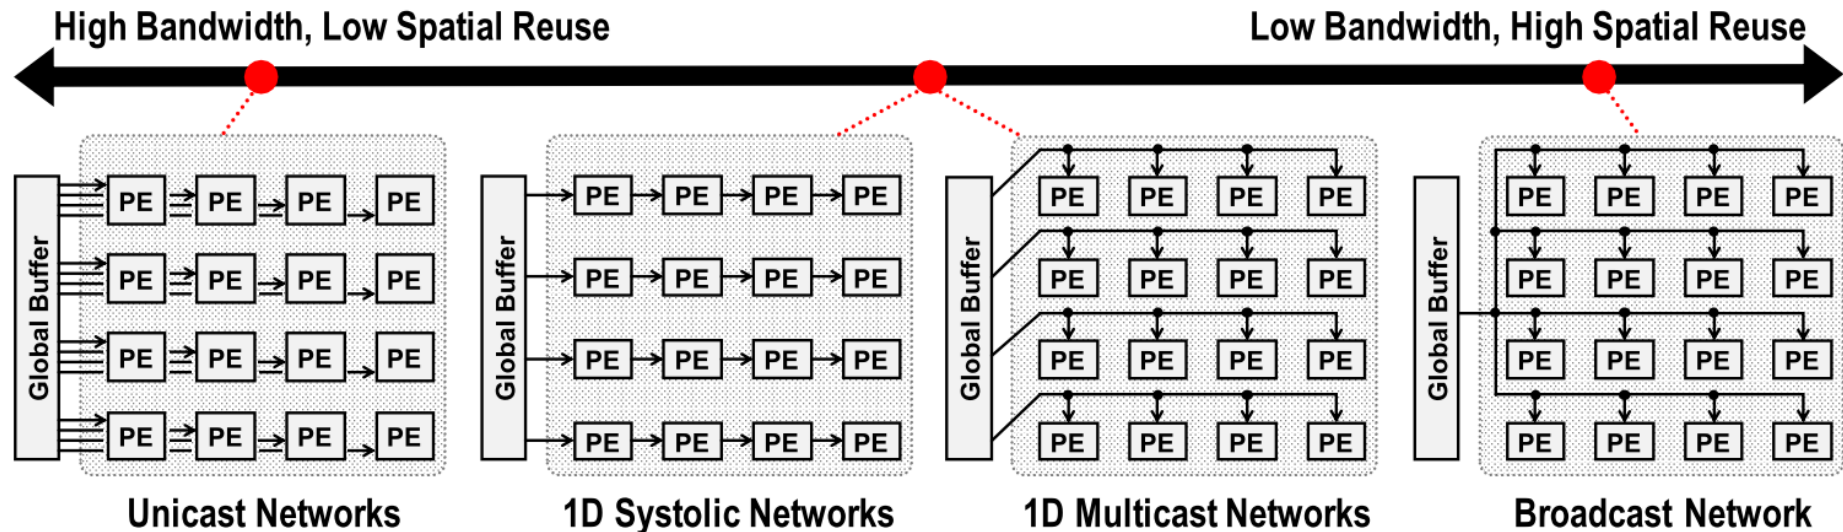
- ❖ Use flexible dataflow (Row Stationary) to exploit reuse in any dimension of DNN to increase energy efficiency and array utilization



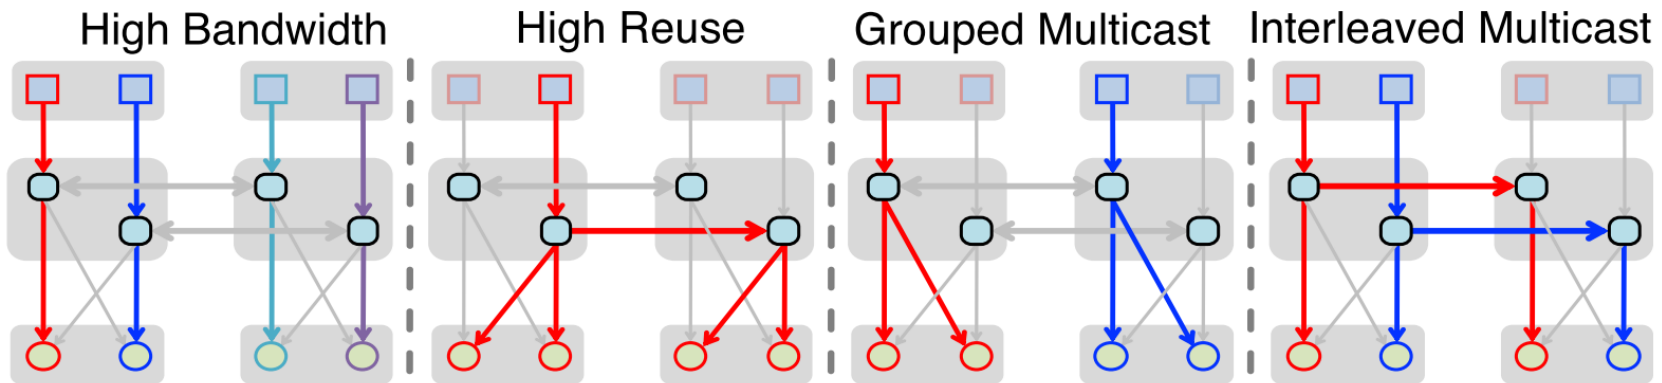

Example: Depth-wise layer

Need Flexible On-Chip Network for Varying Reuse

- ❖ When reuse available, need multicast to exploit spatial data reuse for energy efficiency and high array utilization
- ❖ When reuse not available, need unicast for high BW for weights for FC and weights & activations for high PE utilization
- ❖ An all-to-all on-chip network satisfies above but too expensive and not scalable




Hierarchical Mesh



Eyeriss v2: Balancing Flexibility and Efficiency

Efficiently supports

- ❖ Wide range of filter shapes
 - ❖ Large and Compact
- ❖ Different Layers
 - ❖ CONV, FC, depth wise, etc.
- ❖ Wide range of sparsity
 - ❖ Dense and Sparse
- ❖ Scalable architecture



Speed up over Eyeriss v1 scales with number of PEs

# of PEs	256	1024	16384
AlexNet	17.9x	71.5x	1086.7x
GoogLeNet	10.4x	37.8x	448.8x
MobileNet	15.7x	57.9x	873.0x

Over an order of magnitude faster and more energy efficient than Eyeriss v1

Design Considerations for ASIC

❖ Increase PE utilization

- ❖ Flexible mapping and on-chip network for different DNN models → requires additional hardware

❖ Reduce data movement

- ❖ Custom memory hierarchy and dataflows that exploit data reuse
- ❖ Apply compression to exploit redundancy in data → requires additional hardware

❖ Reduce time and energy per MAC

- ❖ Reduce precision → if precision varies, requires additional hardware; impact on accuracy

❖ Reduce unnecessary MACs

- ❖ Exploit sparsity → requires additional hardware; impact on accuracy
- ❖ Exploit redundant operations → requires additional hardware